

Stanford Open Policing Data - Nashville, TN Analysis

Team: The Brent Spinners

- Kyle Absten
- Ramo Tucakovic
- Ranjita Kaur
- Juan Gutierrez

Introduction

For our final project, we got our hands on a hefty dataset from Stanford University's Open Policing project, specifically traffic stop data from Nashville, Tennessee. We chose the dataset from Nashville TN because it contained a good amount of entries (~3m), with the most complete feature set. Our goal was to build a model that can predict the outcome of these stops: either a warning or something more serious, like a ticket or arrest. This will be a classification solution with only two possible outcomes "WARNING" or "INFRACTION", with infraction encompassing tickets and/or arrests. We want to create a data-driven model to predict the outcome of a traffic stop and investigate which factors are most important to that model.

We begin importing the relevant libraries for our project.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import datetime
import graphviz
import pgeocode
import re
import requests
import time
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder
from sklearn.neighbors import KNeighborsClassifier, KNeighborsRegressor
from sklearn.model_selection import train_test_split, cross_val_score, GridS
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, mean_squared_error
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.inspection import permutation_importance
```

Selection of Data

Source of the Data

The data we chose can be found here: https://stacks.stanford.edu/file/druid:yg821jf8611/yg821jf8611_tn_nashville_2020_04_01.csv.zip

Consider loading the data from a local CSV file, as the download takes several minutes.

```
In [2]: dataset_path = 'https://stacks.stanford.edu/file/druid:yg821jf8611/yg821jf8611_tn_nashville_2020_04_01.csv.zip'
df = pd.read_csv(dataset_path, low_memory=False)
```

Data Exploration

```
In [3]: df.info()
df.sample(3)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3092351 entries, 0 to 3092350
```

```
Data columns (total 42 columns):
```

#	Column	Dtype
0	raw_row_number	object
1	date	object
2	time	object
3	location	object
4	lat	float64
5	lng	float64
6	precinct	object
7	reporting_area	float64
8	zone	object
9	subject_age	float64
10	subject_race	object
11	subject_sex	object
12	officer_id_hash	object
13	type	object
14	violation	object
15	arrest_made	object
16	citation_issued	object
17	warning_issued	object
18	outcome	object
19	contraband_found	object
20	contraband_drugs	object
21	contraband_weapons	object
22	frisk_performed	object
23	search_conducted	object
24	search_person	object
25	search_vehicle	object
26	search_basis	object
27	reason_for_stop	object
28	vehicle_registration_state	object
29	notes	object
30	raw_verbal_warning_issued	object
31	raw_written_warning_issued	object
32	raw_traffic_citation_issued	object
33	raw_misd_state_citation_issued	object
34	raw_suspect_ethnicity	object
35	raw_driver_searched	object
36	raw_passenger_searched	object
37	raw_search_consent	object
38	raw_search_arrest	object
39	raw_search_warrant	bool
40	raw_search_inventory	object
41	raw_search_plain_view	object

```
dtypes: bool(1), float64(4), object(37)
```

```
memory usage: 970.3+ MB
```

Out[3]:

	raw_row_number	date	time	location	lat	lng
2923723	1860668	2014-09-14	01:29:00	MYATT DR & MYATT DR, MADISON, TN, 37115	36.279954	-86.691211
2702709	2215034	2015-08-13	11:33:00	BRILEY PKWY N & MURFREESBORO PIKE, NASHVILLE, ...	36.122074	-86.702419
635941	1555104	2013-12-24	08:17:00	S 7TH ST & SYLVAN ST, NASHVILLE, TN, 37206	36.168169	-86.757240

3 rows × 42 columns

Looking at our dataset, we have 42 different features. 1 feature is boolean, 4 are numeric and the remaining 37 are object. Not all of these features will be helpful to predict whether a future police interaction will result in an infraction or not. A few thoughts on some of the features in our dataset:

raw_row_number- From the documentation: "A number used to join clean data back to the raw data". For our purposes this is not needed.

date- This is the date of the stop. We will look at converting this feature to day_of_week, and see if there is any predictive correlation between the day of the week and whether an infraction is given.

location, lat, long, precinct, reporting_area, zone- These are all variable that deal with where the stop happened and/or the specific police department involved in the stop. It's likely that we will need to drop some of these columns, as they appear to be mostly repetitive data.

officer_id_hash- This is an identifier for the officer involved in the stop. From the documentation

type- From documentation: "Type of stop: vehicular or pedestrian."

violation- From documentation: "Specific violation of stop where provided. What is recorded here varies widely across police departments." Will need to investigate further.

arrest_made, citation_issued, warning_issued, outcome- These are all outcomes of the stop. Because we are only aiming to predict whether the driver will result in an infraction(anything greater than a warning), all of these features can be combined into

one: outcome. This will be our target variable.

contraband_found, contraband_drugs, contraband_weapons- These are all variables related to the subject being found to have illegal items. Need to investigate more to decide whether to include these in the model.

search_conducted, search_person, search_vehicle, search_basis- Similar to contraband data, needs further investigation.

notes- From documentation: "A freeform text field containing any officer notes." Needs further look, but likely needs to be dropped

raw_verbal_warning_issued, raw_written_warning_issued, raw_traffic_citation_issued, raw_misd_state_citation_issued, raw_suspect_ethnicity, raw_driver_searched, raw_passenger_searched, raw_search_consent, raw_search_arrest, raw_search_warrant, raw_search_inventory, raw_search_plain_view- These are all raw versions of other variables in the dataset and can therefore be dropped.

Overall, our dataset is primarily categorical data. Knowing this, we will likely use a decision tree for our model.

Preprocessing and Preparing Data for Modeling

```
In [4]: features_to_drop = ['raw_row_number', 'raw_verbal_warning_issued', 'raw_writ
df.drop(features_to_drop,axis=1, inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3092351 entries, 0 to 3092350
Data columns (total 29 columns):
 #   Column                                Dtype
---  -
 0   date                                 object
 1   time                                 object
 2   location                             object
 3   lat                                 float64
 4   lng                                 float64
 5   precinct                             object
 6   reporting_area                       float64
 7   zone                                 object
 8   subject_age                          float64
 9   subject_race                         object
10  subject_sex                           object
11  officer_id_hash                       object
12  type                                  object
13  violation                             object
14  arrest_made                           object
15  citation_issued                       object
16  warning_issued                        object
17  outcome                               object
18  contraband_found                      object
19  contraband_drugs                      object
20  contraband_weapons                    object
21  frisk_performed                       object
22  search_conducted                      object
23  search_person                         object
24  search_vehicle                        object
25  search_basis                          object
26  reason_for_stop                       object
27  vehicle_registration_state            object
28  notes                                 object
dtypes: float64(4), object(25)
memory usage: 684.2+ MB
```

This cuts our feature set from 42 to 29. Now we will look at the remaining features to assess nan values.

```
In [5]: df.isnull().sum()
```

```
Out[5]: date                0
        time                5467
        location            0
        lat                187106
        lng                187106
        precinct           390222
        reporting_area     332393
        zone               390222
        subject_age        839
        subject_race       1850
        subject_sex        12822
        officer_id_hash    11
        type               0
        violation          8020
        arrest_made        28
        citation_issued    320
        warning_issued     337
        outcome            1935
        contraband_found   2964646
        contraband_drugs   2964646
        contraband_weapons 2964646
        frisk_performed    22
        search_conducted   39
        search_person      43
        search_vehicle     41
        search_basis       2964646
        reason_for_stop    8020
        vehicle_registration_state 31791
        notes              2579713
        dtype: int64
```

What jumps out the most is contraband found, contraband drugs, contraband weapons, search basis and notes are almost entirely nan and therefore should likely be dropped from the dataset.

```
In [6]: df['contraband_found'].isnull().sum()/df.shape[0]
```

```
Out[6]: 0.9587029415483559
```

```
In [7]: df[df['notes'].notnull()]['notes'].head(20)
```

```

Out[7]: 19          -----\n
        26          VIN- WBABK----TET----- \n
        38          DRIVER SIDE TAIL LIGHT NOT FUNCTIONING\n
        40          DUI ARREST\n
        42          EXPIRED TAGS\n
        46          SPEEDING\n
        71          PASSENGER SIDE BRAKE LIGHT OUT\n
        80          HEADLIGHT OUT\n
        84          NO COMPUTER\n
        85          dui\n
        86          SUBJECTS VEHICLE CLOCKED AT -- MPH IN A POSTED...
        116         NO COMPUTER\n
        127         NO BELT\n
        130         TURNING AT INTERSECTION VIOLATION COMMITTED AT...
        136         headlight\n
        143         hl\n
        156         BRAKE LIGHT HAS A SHORT IN IT-DRIVER GOT IT BA...
        182         DUI\n
        189         RAN OFF ROAD. STOPPED TO CHECK FOR DUI. DRIVER...
        198         FORMSTREAM WOULD NOT PULL UP THE MOBILE STOP D...
        Name: notes, dtype: object

```

As we can see, contraband found, contraband drugs, contraband weapons, search basis and notes have 95% nan values. Also the notes column is also very sparse and those that are not nan are too hard to interpret. These columns will be dropped.

```

In [8]: features_to_drop = ['contraband_found', 'contraband_drugs', 'contraband_weap
df.drop(features_to_drop, axis=1, inplace=True)

```

```

In [9]: df.info()
df.isnull().sum()

```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3092351 entries, 0 to 3092350
Data columns (total 24 columns):
 #   Column                                Dtype
---  -
 0   date                                 object
 1   time                                 object
 2   location                             object
 3   lat                                  float64
 4   lng                                  float64
 5   precinct                             object
 6   reporting_area                       float64
 7   zone                                 object
 8   subject_age                          float64
 9   subject_race                         object
10  subject_sex                          object
11  officer_id_hash                       object
12  type                                  object
13  violation                             object
14  arrest_made                           object
15  citation_issued                       object
16  warning_issued                       object
17  outcome                              object
18  frisk_performed                       object
19  search_conducted                      object
20  search_person                         object
21  search_vehicle                        object
22  reason_for_stop                       object
23  vehicle_registration_state            object
dtypes: float64(4), object(20)
memory usage: 566.2+ MB

```

```

Out[9]:  date                0
         time                5467
         location            0
         lat                187106
         lng                187106
         precinct           390222
         reporting_area     332393
         zone               390222
         subject_age         839
         subject_race        1850
         subject_sex         12822
         officer_id_hash      11
         type                0
         violation           8020
         arrest_made         28
         citation_issued     320
         warning_issued      337
         outcome             1935
         frisk_performed     22
         search_conducted    39
         search_person        43
         search_vehicle       41
         reason_for_stop     8020
         vehicle_registration_state 31791
dtype: int64

```

```
In [10]: df['type'].unique()
```

```
Out[10]: array(['vehicular'], dtype=object)
```

All type values are the same: 'vehicular', so this feature can be dropped.

```
In [11]: df.drop('type', axis=1, inplace=True)
```

```
In [12]: df['violation'].value_counts()  
df['violation'].isnull().sum()/df['violation'].shape[0]
```

```
Out[12]: 0.002593496016461262
```

This is a good feature. Less than 1% nan and good categorical groupings.

```
In [13]: df['reason_for_stop'].value_counts()
```

```
Out[13]: reason_for_stop  
moving traffic violation      1546865  
vehicle equipment violation   996282  
safety violation              186139  
registration                  185756  
seatbelt violation            103199  
investigative stop            56489  
parking violation              8483  
child restraint               1118  
Name: count, dtype: int64
```

Looks like 'reason_for_stop' is a duplicate of 'violation'. Dropping.

```
In [14]: df.drop('reason_for_stop', axis=1, inplace=True)
```

```
In [15]: infraction_variables = ['arrest_made', 'citation_issued', 'warning_issued',  
infractions = df[infraction_variables]  
infractions
```

Out[15]:

	arrest_made	citation_issued	warning_issued	outcome
0	False	False	True	warning
1	False	True	False	citation
2	False	False	True	warning
3	False	False	True	warning
4	False	False	True	warning
...
3092346	False	False	True	warning
3092347	False	False	True	warning
3092348	False	False	True	warning
3092349	False	True	False	citation
3092350	False	True	False	citation

3092351 rows × 4 columns

The arrest_made, citation_issued, warning_issued variables are all captured by the outcome variable. Dropping these variables. We may need to do some feature engineering to transform 'outcome' to a boolean infraction_issued variable.

```
In [16]: infractions_to_drop = ['arrest_made', 'citation_issued', 'warning_issued']
df.drop(infractions_to_drop, axis=1, inplace=True)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3092351 entries, 0 to 3092350
Data columns (total 19 columns):
 #   Column                                  Dtype
---  -
 0   date                                   object
 1   time                                   object
 2   location                               object
 3   lat                                    float64
 4   lng                                    float64
 5   precinct                               object
 6   reporting_area                         float64
 7   zone                                   object
 8   subject_age                           float64
 9   subject_race                           object
10  subject_sex                             object
11  officer_id_hash                         object
12  violation                               object
13  outcome                                 object
14  frisk_performed                         object
15  search_conducted                       object
16  search_person                           object
17  search_vehicle                         object
18  vehicle_registration_state             object
dtypes: float64(4), object(15)
memory usage: 448.3+ MB

```

```
In [17]: df.isnull().sum()
```

```

Out[17]: date                0
         time               5467
         location            0
         lat               187106
         lng               187106
         precinct          390222
         reporting_area    332393
         zone              390222
         subject_age        839
         subject_race       1850
         subject_sex       12822
         officer_id_hash     11
         violation          8020
         outcome            1935
         frisk_performed     22
         search_conducted     39
         search_person        43
         search_vehicle       41
         vehicle_registration_state  31791
         dtype: int64

```

Feature Engineering and Advanced Munging

First we will try to convert some of those zip codes from our 'location' feature into usable data to fill in some missing latitude and longitude values.

```
In [18]: # Select rows where 'lat' is null and extract zip from 'location'
no_lat = df[df['lat'].isnull()].copy()
no_lat.loc[:, 'zip_code'] = no_lat['location'].str.extract(r'(\d{5})')
no_lat.head()
no_lat['zip_code'].isnull().sum()
```

Out[18]: 79785

It appears we have ~80k or around 2% of rows with no lat/long data. The first API we found to process this geolocation data limits requests to one a second, meaning that with nearly 80k requests it would take this code a little over twenty-two hours to convert our zip code data into latitudes and longitudes. As an academic exercise and proof of concept we will apply this method to a small sample of our data but the computational and time costs made this first approach impracticable.

```
In [19]: sample = no_lat.sample(10)
sample['lat']

def make_api_request(address):
    url = f"https://geocode.maps.co/search?q={address}&api_key=65d5ac3a7cde5"
    response = requests.get(url)
    if response.status_code == 200:
        data = response.json()
        for result in data:
            if any(keyword in result.get('display_name', '') for keyword in
                    time.sleep(1) # Add delay of 1 second before returning
                return result.get('lat'), result.get('lon') # Safely get la
        print(f"No suitable result found for address: {address}")
        time.sleep(1) # Add delay of 1 second before returning
        return None, None
    else:
        print(f"Failed to retrieve data for address: {address}")
        time.sleep(1) # Add delay of 1 second before returning
        return None, None

# Apply function to 'zip_code' column and assign lat and lng to new columns
sample[['lat', 'lng']] = sample['zip_code'].apply(lambda x: pd.Series(make_a
sample['lat']
```

```
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
No suitable result found for address: nan
```

```

Out[19]: 2552600      None
          302799      None
          1272964    36.16252533066202
          1808288      None
          2359361      None
          2836578      None
          1624489      None
          2708698      None
          153407      None
          2961741      None
          Name: lat, dtype: object

```

Though this first approach would be too slow for our dataset, our second approach using the pgeocode library was able to handle much larger requests:

```

In [20]: def zip_to_lat_long(zip_code):
          nomi = pgeocode.Nominatim('us')
          location = nomi.query_postal_code(zip_code.iloc[0])
          latitude = location['latitude']
          longitude = location['longitude']

          return pd.Series([latitude, longitude])

df['zip_code'] = df['location'].str.extract(r'(\d{5})')
nan_lat_mask = df['lat'].isnull()
zips_to_pull = df.loc[nan_lat_mask, 'zip_code']
zip_codes = zips_to_pull.dropna()
zip_codes.drop_duplicates(inplace=True)
zip_codes.reset_index(drop=True, inplace=True)
zip_codes.columns=['zip_code']
zip_codes = zip_codes.to_frame()
zip_codes[['lat', 'lng']] = zip_codes.apply(lambda x: zip_to_lat_long(x), axis=1)

df.isna().sum()

df.loc[nan_lat_mask, 'lat'] = df[nan_lat_mask].zip_code.map(dict(zip_codes[['zip_code', 'lat']]))
df.loc[nan_lat_mask, 'lng'] = df[nan_lat_mask].zip_code.map(dict(zip_codes[['zip_code', 'lng']]))
df.drop(['location'], inplace=True, axis=1)

```

Next we will address our date/time values to extract some possibly predictive features, such as the day of the week, day of the month, and month.

```

In [21]: df['day_of_week'] = pd.to_datetime(df['date']).dt.day_name()
          df['month'] = pd.to_datetime(df['date']).dt.month
          df['day_of_month'] = pd.to_datetime(df['date']).dt.day

```

And for the final bit of data refinement, our target is warning/infraction, so we are converting arrests and citations to infractions:

```

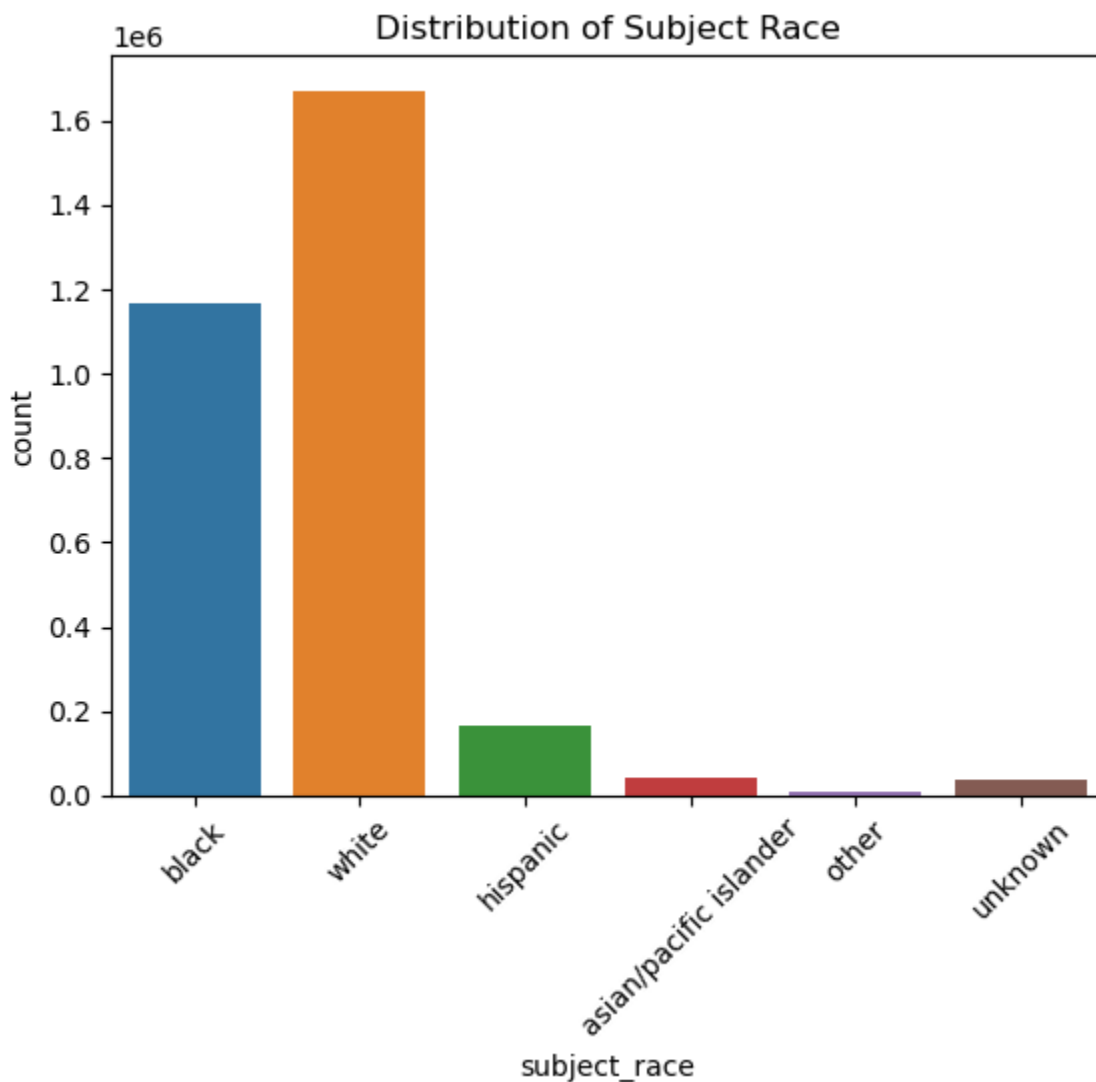
In [22]: df['outcome'] = df['outcome'].apply(lambda x: x if x=='warning' else 'infrac

```

Visualizations

Now that our data has been trimmed down to remove duplicates, most empty values, and probably irrelevant features to the question at hand, we can begin to do some visualizations. Let's begin with a breakdown of frequencies of stops by subject race, which we suspect may be a significant predictor for our purposes.

```
In [23]: sns.countplot(data=df, x='subject_race')
plt.title('Distribution of Subject Race')
plt.xticks(rotation=45) # Rotate x-axis labels if needed
plt.show()
```



Per the US Census Bureau, the racial composition of Tennessee for our three largest values are 59.1% White, 26.8% Black, 10.8% Hispanic or Latino. (<https://www.census.gov/quickfacts/fact/table/nashvilledavidsonmetropolitangovernmentbalancetennessee/PST045222>) When we compare these values to our dataset we get the following

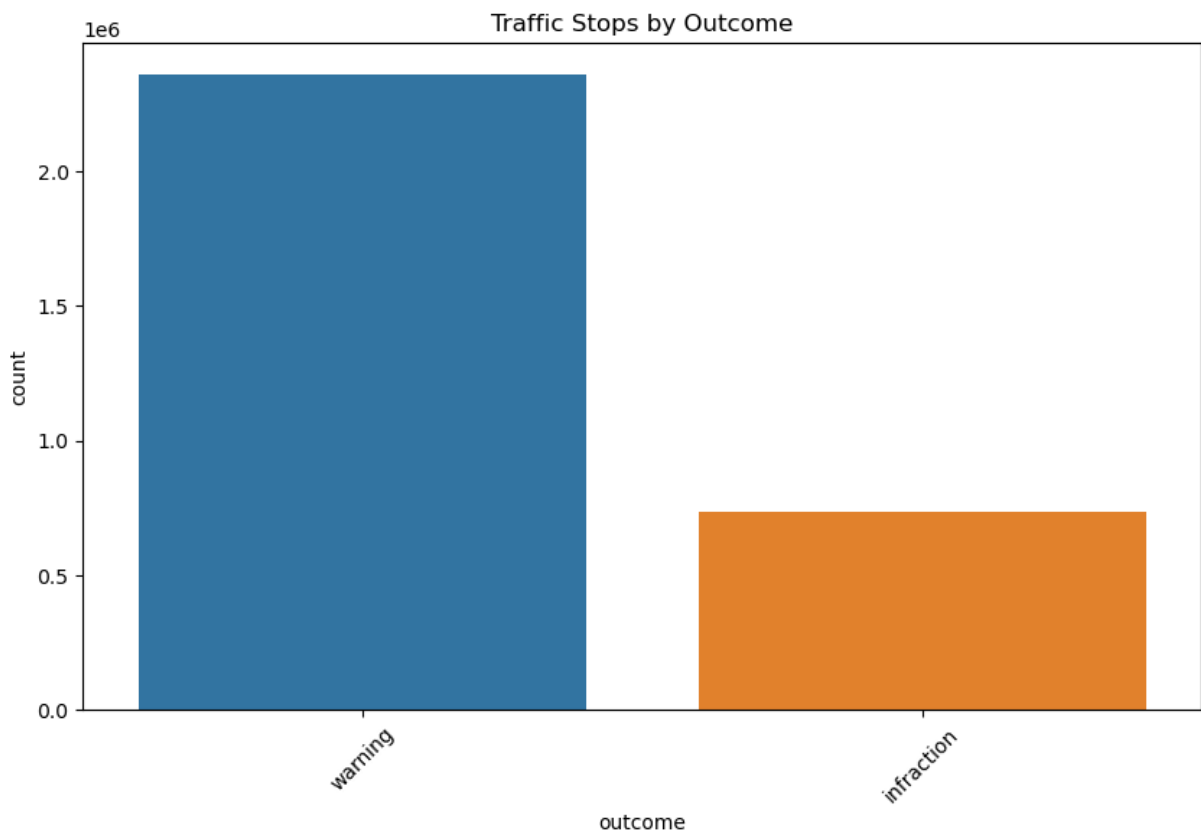
```
In [24]: white_perc = (df[df['subject_race'] == 'white'].shape[0] / df.shape[0])*100
```

```
black_perc = (df[df['subject_race'] == 'black'].shape[0] / df.shape[0])*100
hispanic_perc = (df[df['subject_race'] == 'hispanic'].shape[0] / df.shape[0])
print("White: {:.2f}%, Black: {:.2f}%, Hispanic: {:.2f}%".format(white_perc,
```

White: 54.03%, Black: 37.70%, Hispanic: 5.33%


According to this preliminary data, White people are slightly underrepresented in traffic stops, Black people are significantly overrepresented by around 40% more than expected by demographic breakdown, and Hispanic people are greatly underrepresented, at nearly half the rate of the general population. Next, let's begin taking a look at the outcomes themselves

```
In [25]: plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='outcome', order = df['outcome'].value_counts().ind
plt.title('Traffic Stops by Outcome')
plt.xticks(rotation=45)
plt.show()
```



Distribution by Violation

```
In [26]: plt.figure(figsize=(12, 8))
sns.countplot(data=df, y='violation', order=df['violation'].value_counts().i
plt.title('Distribution of Violations')
plt.xlabel('Count')
plt.ylabel('Violation Type')
plt.show()
```

Ratios of Infractions to Warnings by Race

```
In [27]: outcome_counts = df.groupby('subject_race')['outcome'].value_counts().unstack()
outcome_counts['infraction_to_warning_ratio'] = outcome_counts['infraction'] / outcome_counts['warning']

print("Ratio of infractions to warnings by race:")
print(outcome_counts['infraction_to_warning_ratio'])

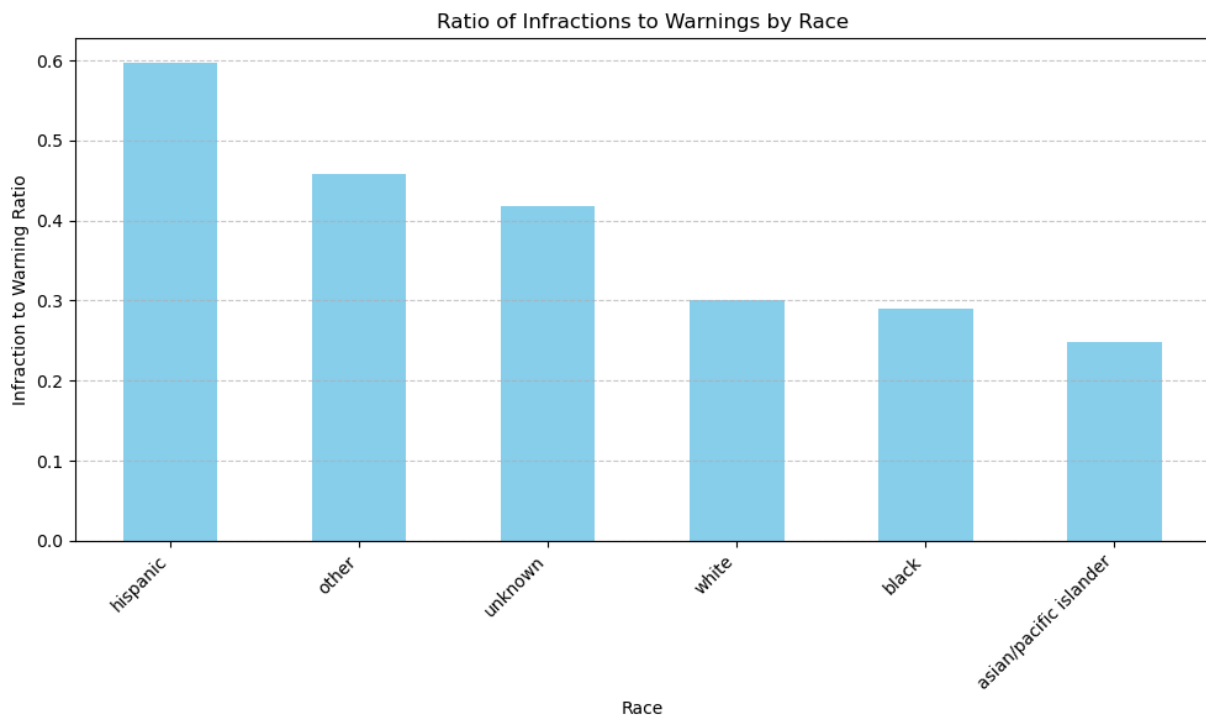
outcome_counts_sorted = outcome_counts.sort_values(by='infraction_to_warning_ratio', ascending=False)

plt.figure(figsize=(10, 6))
outcome_counts_sorted['infraction_to_warning_ratio'].plot(kind='bar', color='red')
plt.title('Ratio of Infractions to Warnings by Race')
plt.xlabel('Race')
plt.ylabel('Infraction to Warning Ratio')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Ratio of infractions to warnings by race:

subject_race	
asian/pacific islander	0.248405
black	0.290060
hispanic	0.597406
other	0.457387
unknown	0.417730
white	0.300546

Name: infraction_to_warning_ratio, dtype: float64



These initial evaluations of our data imply that despite being underrepresented in total number of stops, Hispanic people are more likely to receive a citation or be arrested during a stop.

Methods

Everything that follows was a process of trial and error. Our dataset presented some challenges as we investigated the best model to train for our predictions. In particular the overwhelming amount of categorical features made things difficult. For instance, the `officer_id_hash` feature has 2,295 unique values. If we were to use onehot encoding, this would make our dataset very wide.

As we looked at models, `LogisticRegression` was considered and we cleaned and engineered our data for that type of model. This included one-hot encoding what we could, scaling features and dropping categorical features that didn't make any sense. Our results were barely better than just guessing "warning" for every outcome.

`KNNClassifier` gave slightly better results, however proved to be mostly unusable based on our sample size of 3million+.

We explored the `DecisionTreeClassifier`. Going into the project we figured this would be our model. It deals with categorical data well, so we didn't have to drop as many features, however, it does not deal with nan values. So, after cleaning nan values the best we could, we ended up with 2.7million samples, which was fine. However, we couldn't get the model to do much better than 79%, which is a 2% improvement on our baseline of guessing "warning" for every outcome. This testing was done with

gridsearch and feature sweeping.

After a suggestion from the professor, we looked into HistGradientBoostingClassifier. This model is a bit of a black box for us, and a huge jump from the basic models we learned this semester. However it offers a lot of advantages. It works very well with categorical data, it works around nan values, and it is optimized for speed on large datasets. Without any feature selection or hyper-parameter tuning the model gave us around 85%, which was way better than our previous models.

All of the following data engineering is based around our use of this model.

```
In [28]: df_encoded = df.copy()
df_encoded['precinct'] = pd.to_numeric(df_encoded['precinct'], errors='coerce')
df_encoded['zone'] = pd.to_numeric(df_encoded['zone'], errors='coerce')
```

There were a few typos in the precinct and zone features, so we are converting them to numeric and coercing the errors to nan since our model will handle them just fine.

The HistBoostGradient model works really well with categorical features, so we are using an OrdinalEncoder to convert our categorical features to numerical representations.

```
In [29]: ord_encoder = OrdinalEncoder(handle_unknown="use_encoded_value", unknown_value=-1)

cat_columns = [
    'date',
    'time',
    'precinct',
    'zone',
    'subject_race',
    'subject_sex',
    'officer_id_hash',
    'violation',
    'outcome',
    'frisk_performed',
    'search_conducted',
    'search_person',
    'search_vehicle',
    'vehicle_registration_state',
    'zip_code',
    'day_of_week',
    'month',
    'day_of_month'
]

df_encoded[cat_columns] = ord_encoder.fit_transform(df_encoded[cat_columns])
```

Using GridSearchCV we found the best setting for our model were learning_rate=0.6, max_depth=8. This model is a little confusing for us, but we understand max_depth.

Once we had our best hyper-parameters we did a feature sweep and found that

'day_of_month', 'search_conducted' and 'day_of_week' were noisy features. Our final feature set looks like this:

```
['subject_sex', 'search_vehicle', 'subject_race', 'month', 'zip_code', 'lat', 'officer_id_hash', 'reporting_area', 'search_person', 'date', 'violation', 'time', 'subject_age', 'frisk_performed', 'lng', 'zone', 'precinct', 'vehicle_registration_state']
```

All that is left is to train and test this model

```
In [30]: selected_features = ['subject_sex', 'search_vehicle', 'subject_race', 'month', 'zip_code', 'lat', 'officer_id_hash', 'reporting_area', 'search_person', 'date', 'violation', 'time', 'subject_age', 'frisk_performed', 'lng', 'zone', 'precinct', 'vehicle_registration_state']

X = df_encoded[selected_features]
y = df_encoded['outcome']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

hist_gradient = HistGradientBoostingClassifier(learning_rate=0.6, max_depth=5)
hist_gradient.fit(X_train, y_train)

hist_gradient.score(X_test, y_test)
```

Out[30]: 0.8467510181027179

Results and Discussion

As you can see, we got close to 85% accuracy using this model. This model was a vast improvement over all the other models. However, one down side is that the model is not very interpretable. It is difficult to figure out which features had the largest effect, as is the case for most "black box" models, but luckily the sklearn toolkit includes some testing of permutations of the model to determine which features were most predictive.


```
In [31]: perm_importance = permutation_importance(hist_gradient, X_test, y_test, n_repeats=10, random_state=42)

feature_names = X_test.columns

fi_df = pd.DataFrame({'Feature': feature_names, 'Importance': perm_importance['mean_permutation_importance']})

fi_df = fi_df.sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=fi_df)
plt.title('Permutation Feature Importance')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```



These results were surprising to us, because despite the clear disparities in policing outcomes by race in the data at large and in general, any particular outcome is more strongly predicted by the individual officer who performs the stops, followed by if your traffic stop was escalated to include a search, then by the kind of violation you were stopped for. So to summarize, racial bias in policing is evident from the data as a whole, but the zealousness or permissiveness of particular officers are a much stronger predictor of whether someone gets a ticket or not.